

Pedaling
Towards
Progress

Maxwell
Lindsay

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Pedaling Towards Progress

Analyzing Washington, DC's Bikesharing system
using Open-source tools

Maxwell Lindsay

Van Oord

June 28, 2023



Introduction

My Goals

My Approach

Results

Takeaways

Thanks

About me

- Originally from Washington, DC area
- Background in Coastal Engineering
- Currently work as a Geospatial Developer at Van Oord, In Rotterdam, The Netherlands

Introduction

My Goals

My Approach

Results

Takeaways

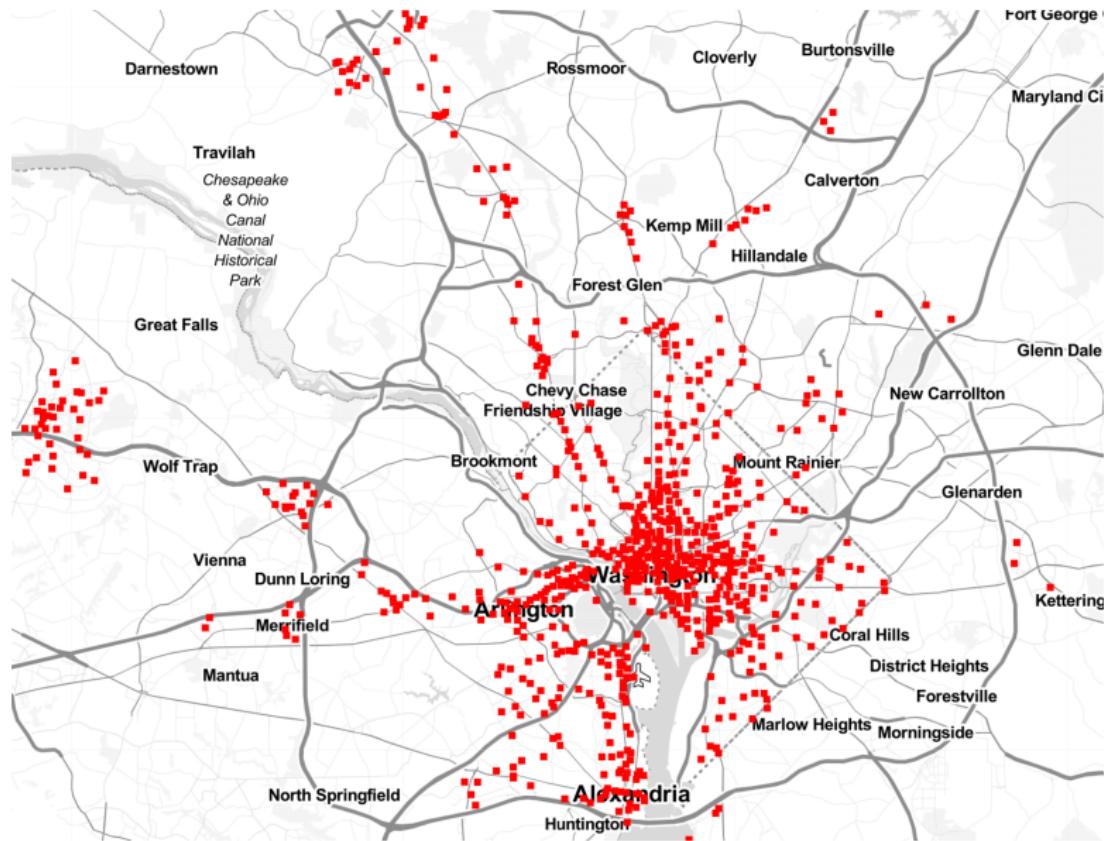
Thanks

Capital Bikeshare basics



- Docked bikeshare
- Primarily intended for short trips
- 700+ docking stations, ~ 35 million trips (so far)

The current station map



Introduction

My Goals

My Approach

Results

Takeaways

Thanks

The data

Capital Bikeshare publishes the following data about each trip on a monthly basis:

- start time
- end time
- start station index
- end station index
- Membership status of the rider (Member or Non-member?)

These data are available as csv files per month or year from Capital Bikeshare at

<https://ride.capitalbikeshare.com/system-data>

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Example data

The raw CSV files are available from the Capital Bikeshare Data portal. This is a small example what the data looks like

ride_id	rideable_type	started_at
5F3D280238A782FE	docked_bike	2023-05-12 18:57:
97EC218DACB24849	classic_bike	2023-05-23 07:55:
31D19AC7BA317018	electric_bike	2023-05-05 17:27:

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Goal 1

Count how many trips occurred on every unique combination of stations

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Goal 2

Find the route of this each of these trips

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Goal 3

Combine these two items to get a map of the estimated number of bikeshare trips on every street/road/path in the city

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

The Challenges

- Efficiently putting hundreds of thousands of trips through a routing engine
- How to combine these hundreds any thousands of complex multiline geometries to sum the trips on each part (without melting my laptop)

Overview of my approach

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

- ① Download all CSV files for all trips
- ② Parse and normalize the CSVs using Pandas
- ③ Find The number of trips between unique pairs of stations
- ④ Build Valhalla routing tiles
- ⑤ Find a route between each station pair using Valhalla
- ⑥ Aggregate the trip statistics across every single route

Data cleaning and validation

For each trip, the trip time is calculated. Trips are considered invalid based on the following criteria:

- Longer than a 4 hours
- Starting and ending at the same station
- Stations with an invalid start or end station

As of May 2023 there are 35,231,413 trips. 175,603 are longer than 4 hours and are removed.

Find unique trips between station pairs

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

35,136,810 trips → 183,959 station combinations

If we drop trips that start and end at the same dock, we are down to 183,214 trips. If we consider only which two stations are involved:

183,214 → 105,636

These trips are the final ones to route

Introduction

My Goals

My Approach

Results

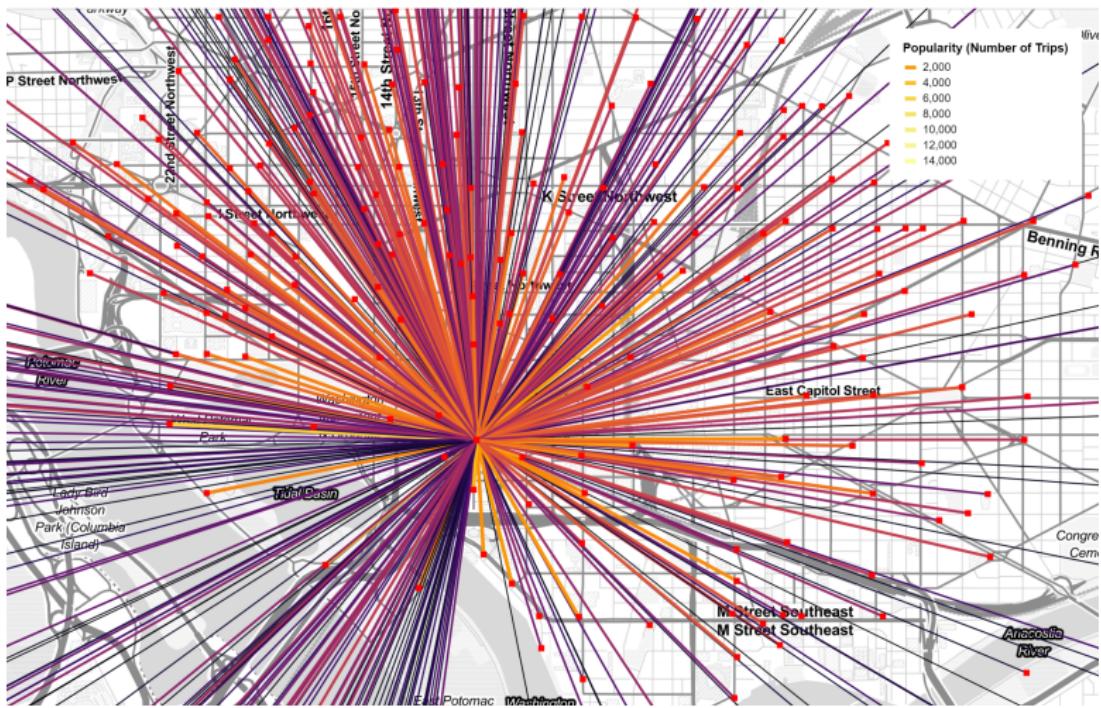
Takeaways

Thanks

Find the number of trips between any 2 stations

Implemented by looping over the individual trips in a Pandas DataFrame.

Results of trips between stations



Building a routing network from OSM data

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

- Download the OSM data export from geoFabrik in protobuf format
- Trim and merge the protobuf OSM files

Routing these trips

Back to Goal 2



Pedaling Towards Progress

Maxwell
Lindsay

Introduction

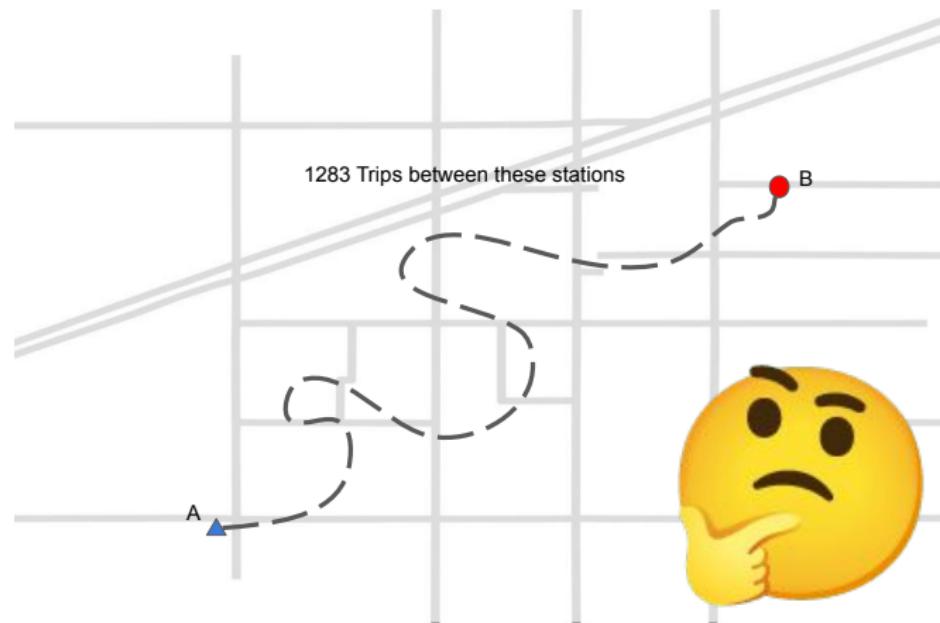
My Goals

My Approach

Results

Takeaways

Thanks



Pedaling
Towards
Progress

Maxwell
Lindsay

Introduction

My Goals

My Approach

Results

Takeaways

Thanks



Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Bike Routing Considerations

Vallhalla allows very granular control over bicycle routing, allowing us to set parameters for a cyclist's:

- Bicycle Type

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Bike Routing Considerations

Vallhalla allows very granular control over bicycle routing, allowing us to set parameters for a cyclist's:

- Bicycle Type
- Willingness to share roads with cars

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Bike Routing Considerations

Vallhalla allows very granular control over bicycle routing, allowing us to set parameters for a cyclist's:

- Bicycle Type
- Willingness to share roads with cars
- Willingness to bike up hills

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Bike Routing Considerations

Vallhalla allows very granular control over bicycle routing, allowing us to set parameters for a cyclist's:

- Bicycle Type
- Willingness to share roads with cars
- Willingness to bike up hills
- Average cycling speed

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Bike Routing Considerations

Vallhalla allows very granular control over bicycle routing, allowing us to set parameters for a cyclist's:

- Bicycle Type
- Willingness to share roads with cars
- Willingness to bike up hills
- Average cycling speed

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Bike Routing Considerations

Vallhalla allows very granular control over bicycle routing, allowing us to set parameters for a cyclist's:

- Bicycle Type
- Willingness to share roads with cars
- Willingness to bike up hills
- Average cycling speed

How to do this efficiently?

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Bike Routing Considerations

Vallhalla allows very granular control over bicycle routing, allowing us to set parameters for a cyclist's:

- Bicycle Type
- Willingness to share roads with cars
- Willingness to bike up hills
- Average cycling speed

How to do this efficiently?

Avoid the Valhalla HTTP API and use *pyvalhalla*

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Back to Goal 3

How to combine it all

The challenge:

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Back to Goal 3

How to combine it all

The challenge:

How to join over 100k linestrings, summing a certain value *only where they overlap*

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Back to Goal 3

How to combine it all

The challenge:

How to join over 100k linestrings, summing a certain value *only where they overlap*

The solution I found:

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Back to Goal 3

How to combine it all

The challenge:

How to join over 100k linestrings, summing a certain value *only where they overlap*

The solution I found:

TopoGeometry

Pedaling Towards Progress

Maxwell
Lindsay

Introduction

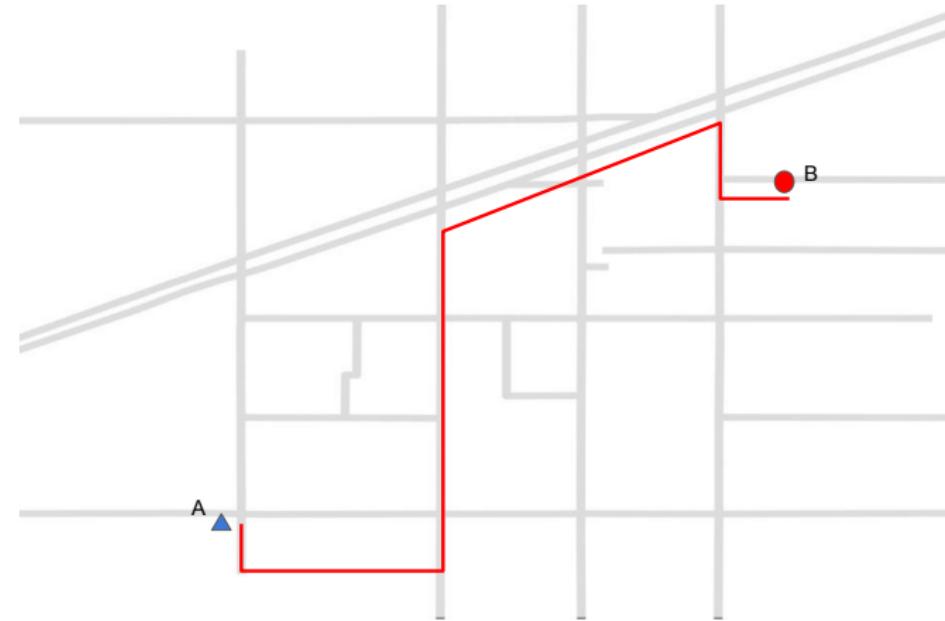
My Goals

My Approach

Results

Takeaways

Thanks



Pedaling Towards Progress

Maxwell
Lindsay

Introduction

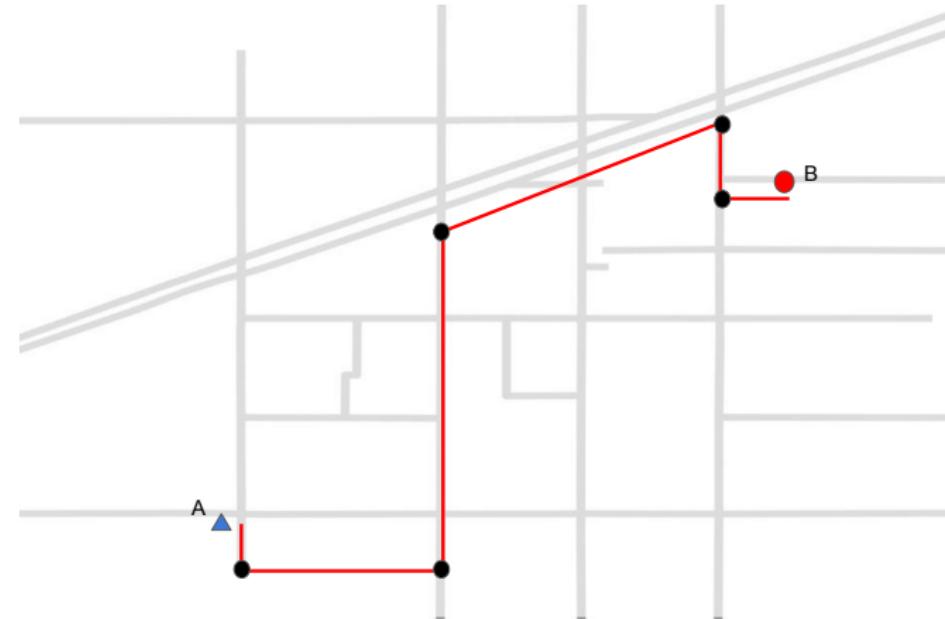
My Goals

My Approach

Results

Takeaways

Thanks



Pedaling Towards Progress

Maxwell
Lindsay

Introduction

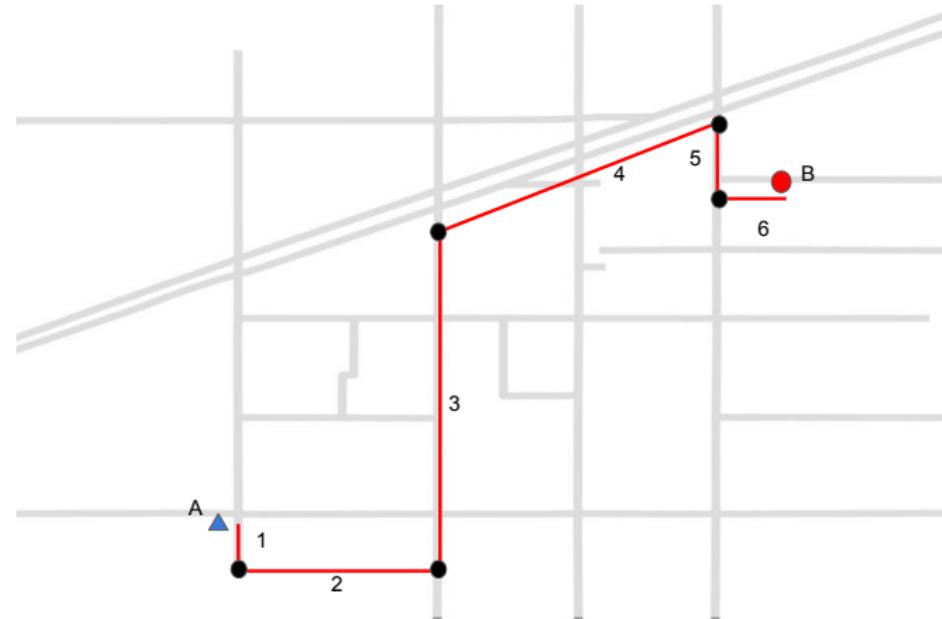
My Goals

My Approach

Results

Takeaways

Thanks



Introduction

My Goals

My Approach

Results

Takeaways

Thanks

PostGIS topology

- Using PLpgsql triggers, add a corresponding entry to the topogeometry table for every new entry to the routes tables

Pedaling Towards Progress

Maxwell
Lindsay

Introduction

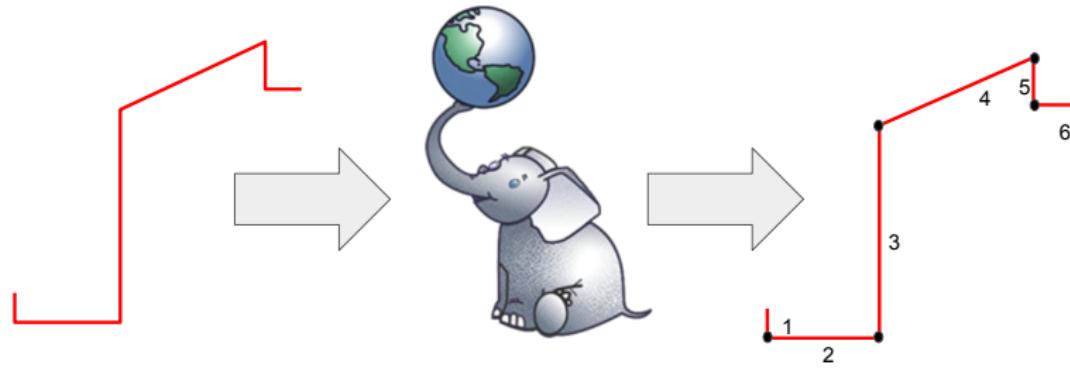
My Goals

My Approach

Results

Takeaways

Thanks



Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Sum up the trips on every topological edge

- Using an SQL query, we can sum the trips on every unique “section” of road in DC

Pedaling Towards Progress

Maxwell
Lindsay

Introduction

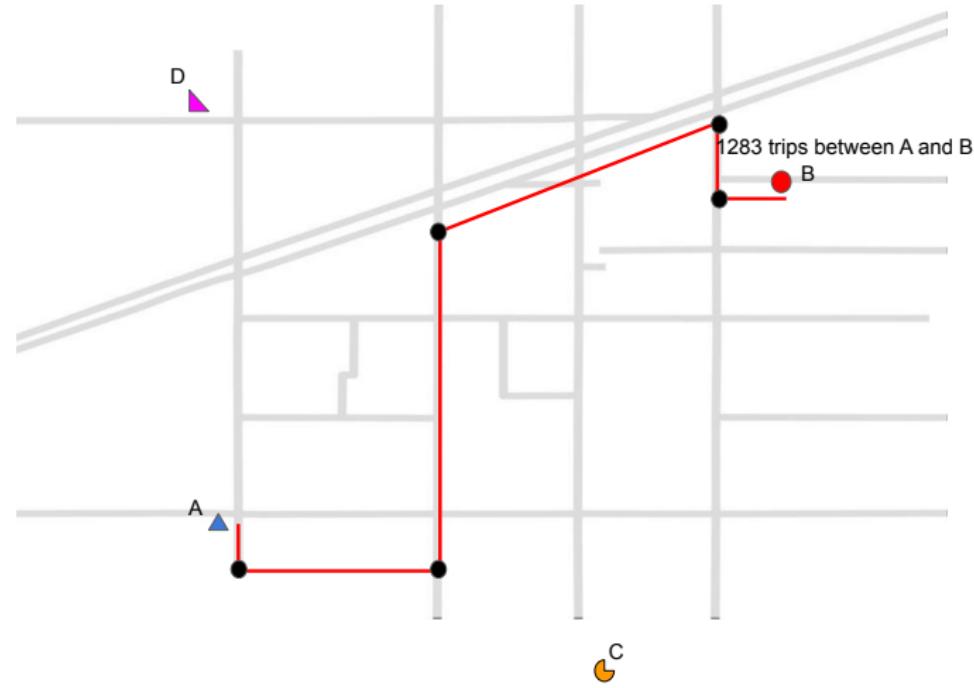
My Goals

My Approach

Results

Takeaways

Thanks



Pedaling Towards Progress

Maxwell
Lindsay

Introduction

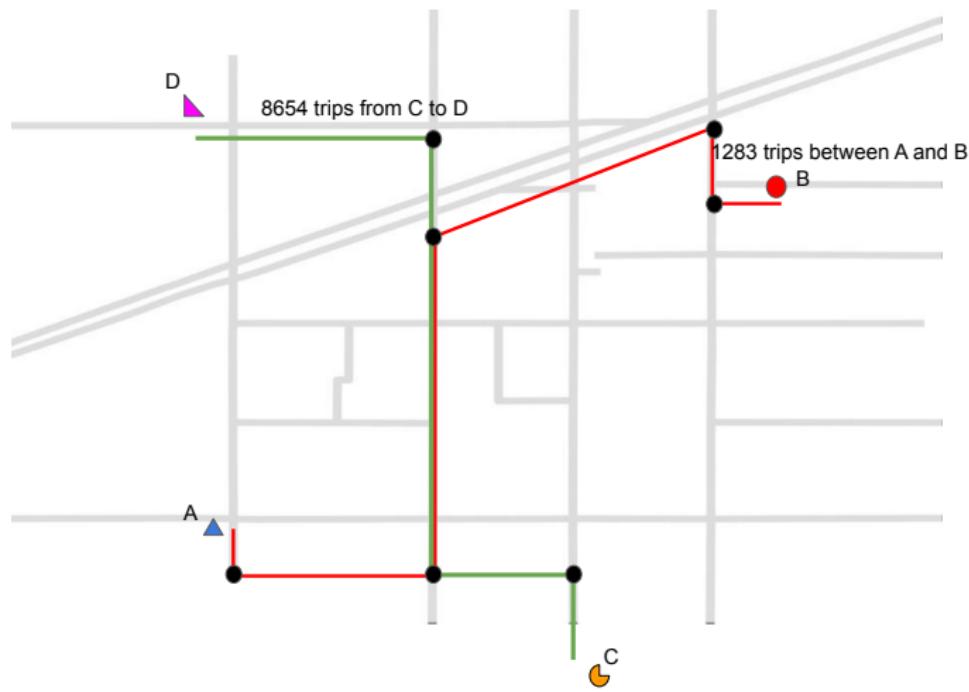
My Goals

My Approach

Results

Takeaways

Thanks



Pedaling Towards Progress

Maxwell
Lindsay

Introduction

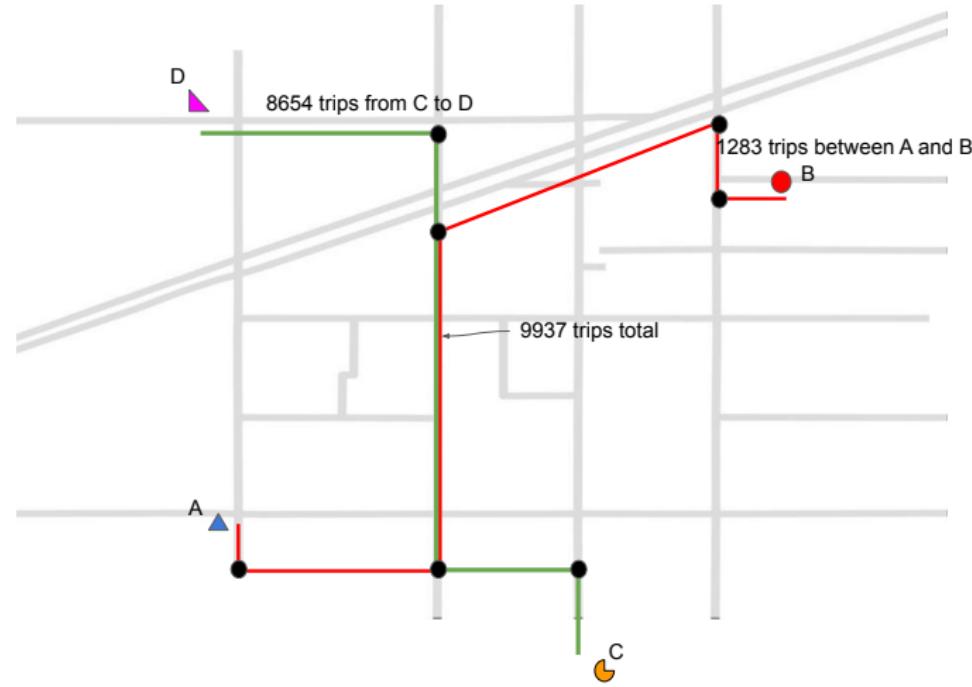
My Goals

My Approach

Results

Takeaways

Thanks



Introduction

My Goals

My Approach

Results

Takeaways

Thanks

The actual implementation

- PostGIS hosted inside a Docker container

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

The actual implementation

- PostGIS hosted inside a Docker container
- A series of Python scripts load the data, and run the routing software

The actual implementation

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

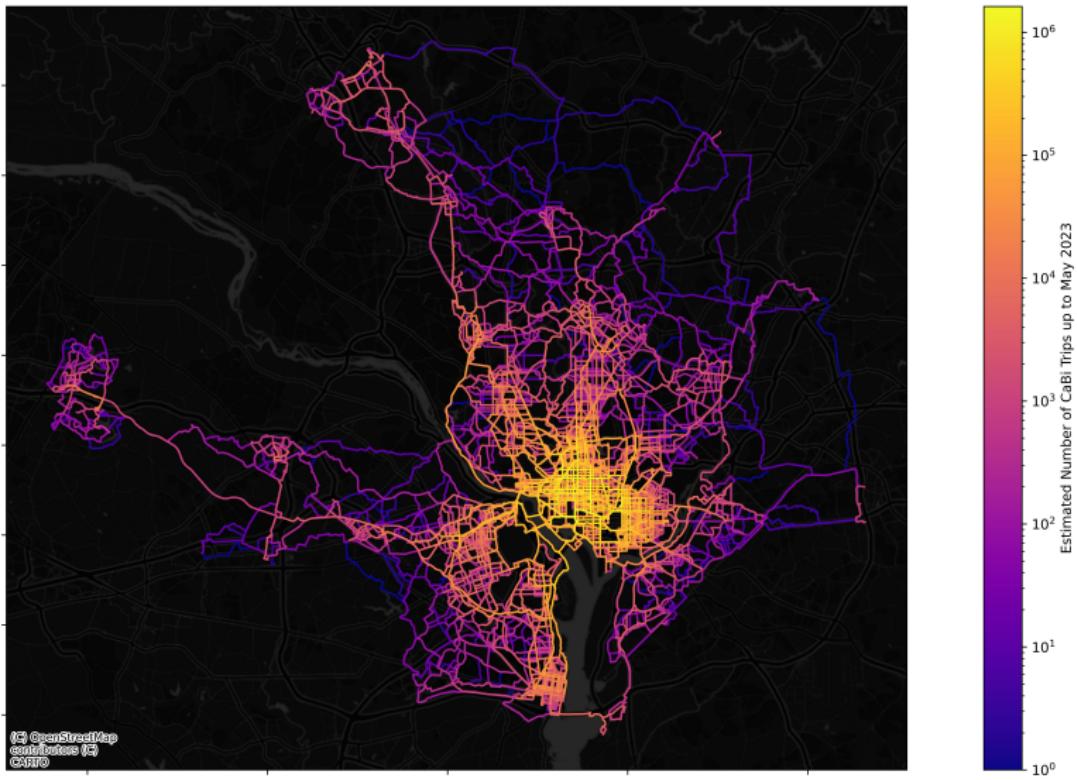
- PostGIS hosted inside a Docker container
- A series of Python scripts load the data, and run the routing software
- Makefile ties everything together

Pedaling Towards Progress

Maxwell
Lindsay

Introduction
My Goals
My Approach
Results
Takeaways
Thanks

Results



Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Validation

How accurate is this?

Are the sums for any individual routes anywhere near accurate?

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Validation

How accurate is this?

Are the sums for any individual routes anywhere near accurate?

Probably not

Which steps made this practical to run on a laptop?

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

- using *pyvalhalla* instead of hosting it in Docker
- Using topology to sum trips by street
- Using SQL triggers for creating the topogeometry

Introduction

My Goals

My Approach

Results

Takeaways

Thanks

What would I have done differently?

- more in SQL, less in Python
- Python is perfect for being the "Glue" and interacting with the various parts

Introduction

My Goals

My Approach

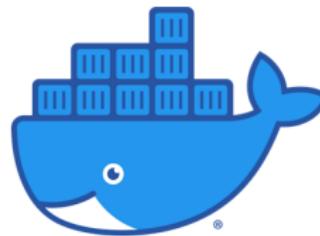
Results

Takeaways

Thanks

Shoutouts

PostGIS



GeoPandas



Introduction

My Goals

My Approach

Results

Takeaways

Thanks

Thanks for your interest